

## CSE 655 – Project Proposal

Student: Merve DÖNMEZ (244201001016)

### Project Topic

Privacy-Preserving Credit Default Prediction using Homomorphic Encryption (SEAL CKKS) and Stacking Ensemble

### Motivation

Financial datasets contain sensitive personal and transactional information. Organizations often face regulatory and confidentiality constraints that limit data sharing and model deployment across parties. This project explores privacy-preserving machine learning using approximate homomorphic encryption (CKKS) to protect data during processing while maintaining competitive predictive performance. We target credit default prediction as a representative and high-impact use case and evaluate the trade-offs between privacy, accuracy, runtime, and memory.

Key motivations:

- Maintain data confidentiality during model training/inference via homomorphic encryption.
- Quantify the overhead introduced by encryption and identify practical optimization strategies (batching, parameter tuning).
- Provide a reproducible pipeline combining strong tabular models through stacking with HE-based data handling.

### Available Datasets

- UCI Default of Credit Card Clients dataset (UCI repository, ucimlrepo id=350). The dataset includes demographic and payment history attributes to predict credit default probability.
  - Usage plan: feature engineering (ratios, interactions, summary stats), class rebalancing with SMOTE, scaling with MinMax, train/test split.
- Optional alternatives for robustness checks (time permitting): UCI German Credit, public Lending Club snapshots.

Links/references:

- UCI Default of Credit Card Clients (ucimlrepo id=350)

### Relation to Your Graduate Work

The project aligns with privacy-preserving and trustworthy AI research directions, focusing on encrypted computation and secure analytics. It contributes empirical evidence on the feasibility of CKKS-based encrypted pipelines for

tabular credit risk modeling and informs thesis-level exploration of trade-offs (accuracy vs. compute/memory vs. privacy), parameter tuning (scale, modulus, poly degree), and deployment considerations.

## Methodology and Planned System

### 1) Data Pipeline

- Feature engineering: ratio, product, difference, row-wise sums/means/std; feature selection with RandomForest importances (top-10).
- Imbalance handling with SMOTE; MinMax scaling; train/test split.

### 2) Model

- Stacking ensemble with base learners: RandomForest, XGBoost, LightGBM; meta-learner: LogisticRegression.
- Parallelism via `n_jobs=-1` where supported; optional GPU for gradient-boosting (time permitting).

### 3) Privacy (Homomorphic Encryption)

- Microsoft SEAL (CKKS) for approximate arithmetic on encrypted real-valued vectors.
- Parameters (initial): `poly_modulus_degree=4096`; `CoeffModulus.Create(4096, [40,20,40])`; `scale=235`.
- Batch encrypt/decrypt utilities; batch size tuning (initial default 256) to balance throughput and memory.

### 4) Evaluation

- Metrics: ROC-AUC, accuracy, classification report; optimal probability threshold via `argmax(tpr - fpr)` on ROC.
- Performance: measure runtime and peak memory for plain vs. encrypted flows to quantify overhead.
- Ablations (time permitting): batch size (128/256/512), CKKS parameters, early stopping for boosters, float32 vs float64.

## Hardware Requirements

- CPU with multi-core support (recommended 8+ cores). HE operations are CPU-intensive.
- Memory: 16 GB RAM recommended to comfortably handle batch-based encryption and model training.
- Optional GPU (e.g., NVIDIA CUDA) for accelerated XGBoost/LightGBM training; not strictly required.

## Related Literature

- Cheon, Jung Hee, et al. “CKKS: An Approximate Homomorphic Encryption Scheme for Real Numbers.” (Original CKKS scheme introducing

approximate arithmetic on ciphertexts.)

- Microsoft SEAL Documentation. Homomorphic encryption library with CKKS and BFV schemes and implementation best practices.
- CryptoNets and follow-up works on privacy-preserving ML using HE for inference/training, demonstrating feasibility and trade-offs in practical settings.
- Stacking ensemble literature (Wolpert’s stacking and subsequent applications to tabular ML) showing performance gains from meta-learning over diverse base learners.

Please include formal citations in the final PDF (IEEE/ACM/APA) as required by course policy.

### **Project Plan and Milestones**

- Week 1–2: Data ingestion, feature engineering, baseline stacking (plain). Establish metrics.
- Week 3: Integrate CKKS pipeline (encrypt/decrypt), parameterize batch size, measure overhead.
- Week 4: Optimization passes (batch size sweep, early stopping, float32), ablation runs, documentation.

### **Expected Outcomes**

- A working privacy-preserving credit default prediction pipeline with reproducible experiments.
- Quantitative comparison of plain vs. encrypted execution (accuracy, AUC, runtime, memory).
- Recommendations for practical CKKS configurations and batching strategies in tabular credit risk tasks.

### **Repository**

- Source code and experiment artifacts will be maintained at: <https://github.com/MerveDnmz/SealCreditDefaultClassifierWithStackingEnsemble>

### **Submission**

- Deliverable: 2–3 page PDF with the sections above, figures/tables as needed, submitted by October 14, 2025, 11:59 PM.

### **References**

- Cheon, J. H., Kim, A., Kim, M., & Song, Y. (2017). Homomorphic encryption for arithmetic of approximate numbers. In ASIACRYPT 2017 (pp. 409–437). Springer. [https://doi.org/10.1007/978-3-319-70694-8\\_15](https://doi.org/10.1007/978-3-319-70694-8_15)

- Microsoft SEAL (release v4.x) documentation. (n.d.). <https://github.com/microsoft/SEAL> and <https://github.com/microsoft/SEAL/tree/main/native/src/seal>
- Gilad-Bachrach, R., Dowlin, N., Laine, K., Lauter, K., Naehrig, M., & Wernsing, J. (2016). CryptoNets: Applying neural networks to encrypted data with high throughput and accuracy. In ICML 2016 (pp. 201–210). <http://proceedings.mlr.press/v48/gilad-bachrach16.html>
- Wolpert, D. H. (1992). Stacked generalization. *Neural Networks*, 5(2), 241–259. [https://doi.org/10.1016/S0893-6080\(05\)80023-1](https://doi.org/10.1016/S0893-6080(05)80023-1)
- Breiman, L. (1996). Stacked regressions. *Machine Learning*, 24(1), 49–64. <https://doi.org/10.1007/BF00117832>
- Chen, T., & Guestrin, C. (2016). XGBoost: A scalable tree boosting system. In KDD 2016 (pp. 785–794). <https://doi.org/10.1145/2939672.2939785>
- Ke, G., et al. (2017). LightGBM: A highly efficient gradient boosting decision tree. In NeurIPS 2017. [https://papers.nips.cc/paper\\_files/paper/2017/hash/6449f44a102fde848669bdd9eb6b76fa-Abstract.html](https://papers.nips.cc/paper_files/paper/2017/hash/6449f44a102fde848669bdd9eb6b76fa-Abstract.html)